



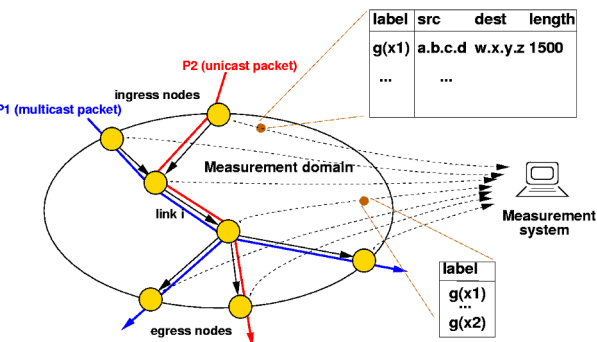


# Sampling for Big Data

Graham Cormode, University of Warwick

[G.Cormode@warwick.ac.uk](mailto:G.Cormode@warwick.ac.uk)

Nick Duffield, Texas A&M University

[Nick.Duffield@gmail.com](mailto:Nick.Duffield@gmail.com)



## Big Data

◇ “Big” data arises in many forms:

- **Physical Measurements**: from science (physics, astronomy)
- **Medical data**: genetic sequences, detailed time series
- **Activity data**: GPS location, social network activity
- **Business data**: customer behavior tracking at fine detail


◇ **Common themes**:

- Data is large, and growing
- There are important patterns and trends in the data
- We don't fully know where to look or how to find them




## Why Reduce?

- ◇ Although “big” data is about more than just the volume...  
...most big data is big!
- ◇ It is not always possible to store the data in full
  - Many applications (telecoms, ISPs, search engines) can’t keep everything
- ◇ It is inconvenient to work with data in full
  - Just because we can, doesn’t mean we should
- ◇ It is faster to work with a compact summary
  - Better to explore data on a laptop than a cluster





## Why Sample?

- ◇ Sampling has an intuitive semantics
  - We obtain a smaller data set with the same structure
- ◇ Estimating on a sample is often straightforward
  - Run the analysis on the sample that you would on the full data
  - Some rescaling/reweighting may be necessary
- ◇ Sampling is general and agnostic to the analysis to be done
  - Other summary methods only work for certain computations
  - Though sampling can be tuned to optimize some criteria
- ◇ Sampling is (usually) easy to understand
  - So prevalent that we have an intuition about sampling



## Alternatives to Sampling

- ◇ Sampling is not the only game in town
  - Many other data reduction techniques by many names
- ◇ Dimensionality reduction methods
  - PCA, SVD, eigenvalue/eigenvector decompositions
  - Costly and slow to perform on big data
- ◇ “Sketching” techniques for streams of data
  - Hash based summaries via random projections
  - Complex to understand and limited in function
- ◇ Other transform/dictionary based summarization methods
  - Wavelets, Fourier Transform, DCT, Histograms
  - Not incrementally updatable, high overhead
- ◇ All worthy of study – in other tutorials



## Health Warning: contains probabilities

- ◇ Will avoid detailed probability calculations, aim to give high level descriptions and intuition
- ◇ But some probability basics are assumed
  - Concepts of probability, expectation, variance of random variables
  - Allude to concentration of measure (Exponential/Chernoff bounds)
- ◇ Feel free to ask questions about technical details along the way

$$\begin{aligned}\text{var} \left( \frac{k}{n} \right) &= \text{E} \left[ \text{var} \left( \frac{k}{n} \middle| \theta \right) \right] + \text{var} \left[ \text{E} \left( \frac{k}{n} \middle| \theta \right) \right] \\ &= \text{E} \left[ \left( \frac{1}{n} \right) \theta(1 - \theta) \middle| \mu, M \right] + \text{var} (\theta | \mu, M) \\ &= \frac{1}{n} (\mu(1 - \mu)) + \frac{n-1}{n} \frac{(\mu(1 - \mu))}{M+1} \\ &= \frac{\mu(1 - \mu)}{n} \left( 1 + \frac{n-1}{M+1} \right).\end{aligned}$$

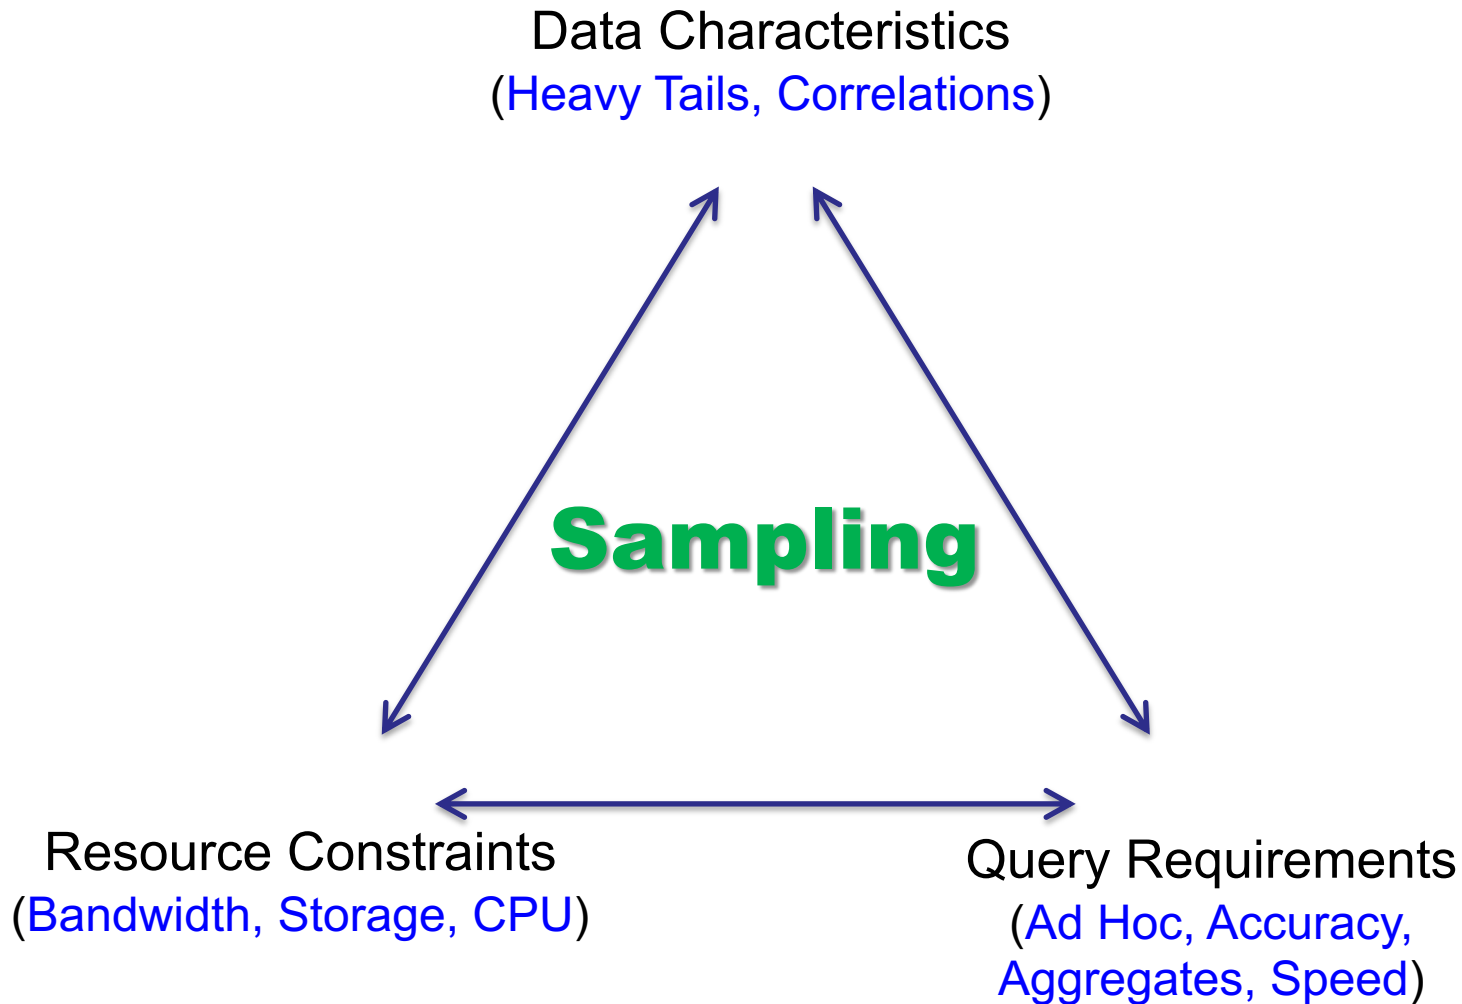
## Outline

- ◇ Motivating application: sampling in large ISP networks
- ◇ Basics of sampling: concepts and estimation
- ◇ Stream sampling: uniform and weighted case
  - Variations: Concise sampling, sample and hold, sketch guided

### BREAK

- ◇ Advanced stream sampling: sampling as cost optimization
  - VarOpt, priority, structure aware, and stable sampling
- ◇ Hashing and coordination
  - Bottom-k, consistent sampling and sketch-based sampling
- ◇ Graph sampling
  - Node, edge and subgraph sampling
- ◇ Conclusion and future directions

# Sampling as a Mediator of Constraints






## Motivating Application: ISP Data

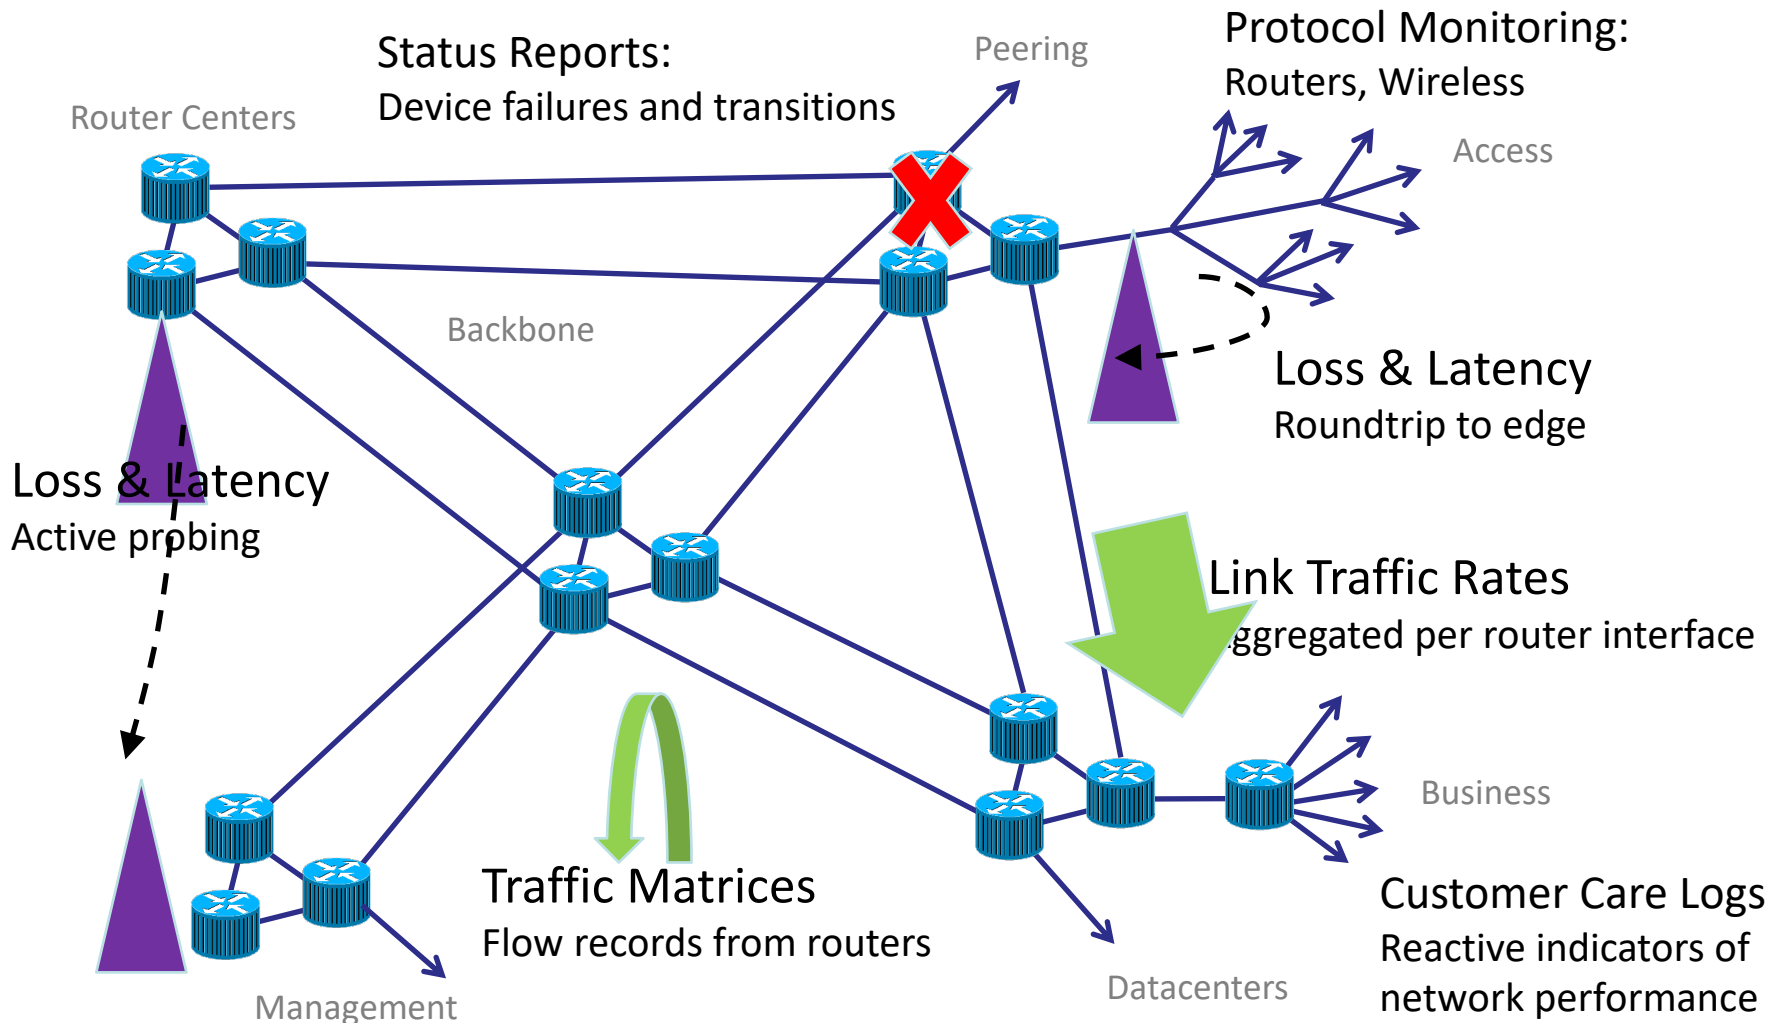
- ◇ Will motivate many results with application to ISPs
- ◇ Many reasons to use such examples:
  - **Expertise**: tutors from telecoms world
  - **Demand**: many sampling methods developed in response to ISP needs
  - **Practice**: sampling widely used in ISP monitoring, built into routers
  - **Prescience**: ISPs were first to hit many “big data” problems
  - **Variety**: many different places where sampling is needed
- ◇ First, a crash-course on ISP networks...



## Structure of Large ISP Networks



# Measuring the ISP Network: Data Sources



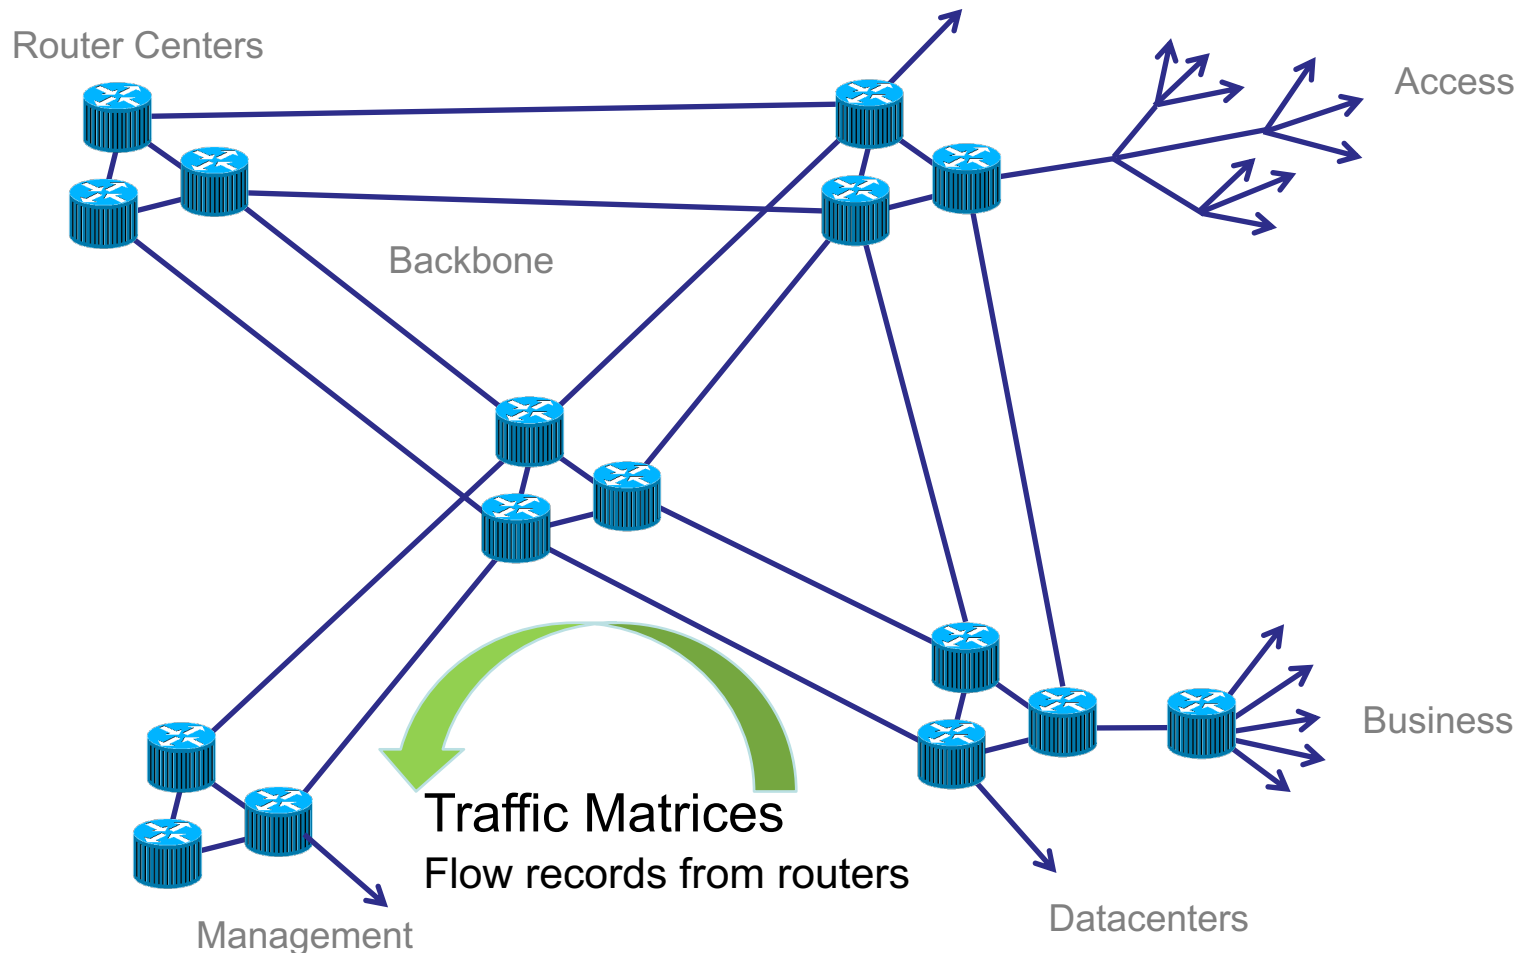
# Why Summarize (ISP) Big Data?

- ◇ When transmission bandwidth for measurements is limited
  - Not such a big issue in ISPs with in-band collection
- ◇ Typically raw accumulation is not feasible (even for nation states)
  - High rate streaming data
  - Maintain historical summaries for baselining, time series analysis
- ◇ To facilitate fast queries
  - When infeasible to run exploratory queries over full data
- ◇ As part of hierarchical query infrastructure:
  - Maintain full data over limited duration window
  - Drill down into full data through one or more layers of summarization


Sampling has been proved to be a flexible method to accomplish this

## **Data Scale: Summarization and Sampling**

# Traffic Measurement in the ISP Network



## Massive Dataset: Flow Records



- ◇ **IP Flow**: set of packets with common key observed close in time
- ◇ **Flow Key**: IP src/dst address, TCP/UDP ports, ToS,... [64 to 104+ bits]
- ◇ **Flow Records**:
  - Protocol level summaries of flows, compiled and exported by routers
  - Flow key, packet and byte counts, first/last packet time, some router state
  - Realizations: Cisco Netflow, IETF Standards
- ◇ **Scale**: 100's TeraBytes of flow records daily are generated in a large ISP
- ◇ Used to manage network over range of timescales:
  - Capacity planning (months),..., detecting network attacks (seconds)
- ◇ Analysis tasks
  - **Easy**: timeseries of predetermined aggregates (e.g. address prefixes)
  - **Hard**: fast queries over exploratory selectors, history, communications subgraphs

# Flows, Flow Records and Sampling

- ◇ Two types of sampling used in practice for internet traffic:
  1. Sampling packet stream in router prior to forming flow records
    - Limits the rate of lookups of packet key in flow cache
    - Realized as Packet Sampled NetFlow (more later...)
  2. Downstream sampling of flow records in collection infrastructure
    - Limits transmission bandwidth, storage requirements
    - Realized in ISP measurement collection infrastructure (more later...)
- ◇ Two cases illustrative of general property
  - Different underlying distributions require different sample designs
  - Statistical optimality sometimes limited by implementation constraints
    - Availability of router storage, processing cycles



## Abstraction: Keyed Data Streams

- ◇ **Data Model**: objects are keyed weights
  - Objects  $(x,k)$ : Weight  $x$ ; key  $k$ 
    - **Example 1**: objects = packets,  $x$  = bytes,  $k$  = key (source/destination)
    - **Example 2**: objects = flows,  $x$  = packets or bytes,  $k$  = key
    - **Example 3**: objects = account updates,  $x$  = credit/debit,  $k$  = account ID
- ◇ Stream of keyed weights,  $\{(x_i, k_i): i = 1, 2, \dots, n\}$
- ◇ Generic query: subset sums
  - $X(S) = \sum_{i \in S} x_i$  for  $S \subset \{1, 2, \dots, n\}$  i.e. total weight of index subset  $S$
  - Typically  $S = S(K) = \{i: k_i \in K\}$  : objects with keys in  $K$ 
    - **Example 1, 2**:  $X(S(K))$  = total bytes to given IP dest address / UDP port
    - **Example 3**:  $X(S(K))$  = total balance change over set of accounts
- ◇ **Aim**: Compute fixed size summary of stream that can be used to estimate arbitrary subset sums with known error bounds

## Inclusion Sampling and Estimation

### ◇ Horvitz-Thompson Estimation:

- Object of size  $x_i$  sampled with probability  $p_i$
- Unbiased estimate  $x'_i = x_i / p_i$  (if sampled), 0 if not sampled:  $E[x'_i] = x_i$

### ◇ Linearity:

- Estimate of subset sum = sum of matching estimates
- Subset sum  $X(S) = \sum_{i \in S} x_i$  is estimated by  $X'(S) = \sum_{i \in S} x'_i$

### ◇ Accuracy:

- Exponential Bounds:  $\Pr[|X'(S) - X(S)| > \delta X(S)] \leq \exp[-g(\delta)X(S)]$
- Confidence intervals:  $X(S) \in [X^-(\varepsilon), X^+(\varepsilon)]$  with probability  $1 - \varepsilon$

### ◇ Futureproof:

- Don't need to know queries at time of sampling
  - “Where/where did that suspicious UDP port first become so active?”
  - “Which is the most active IP address within than anomalous subnet?”
- Retrospective estimate: subset sum over relevant keyset

## Independent Stream Sampling

### ◇ Bernoulli Sampling

- IID sampling of objects with some probability  $p$
- Sampled weight  $x$  has HT estimate  $x/p$

### ◇ Poisson Sampling

- Weight  $x_i$  sampled with probability  $p_i$  ; HT estimate  $x_i / p_i$

### ◇ When to use Poisson vs. Bernoulli sampling?

- Elephants and mice: Poisson allows probability to depend on weight...

### ◇ What is best choice of probabilities for given stream $\{x_i\}$ ?



## Bernoulli Sampling

- ◇ The easiest possible case of sampling: all weights are 1
  - $N$  objects, and want to sample  $k$  from them uniformly
  - Each possible subset of  $k$  should be equally likely
- ◇ Uniformly sample an index from  $N$  (without replacement)  $k$  times
  - Some subtleties: truly random numbers from  $[1...N]$  on a computer?
  - Assume that random number generators are good enough
- ◇ Common trick in DB: assign a random number to each item and sort
  - Costly if  $N$  is very big, but so is random access
- ◇ Interesting problem: take a single linear scan of data to draw sample
  - Streaming model of computation: see each element once
  - **Application**: IP flow sampling, too many (for us) to store
  - (For a while) common tech interview question

# Reservoir Sampling


“Reservoir sampling” described by [Knuth 69, 81]; enhancements [Vitter 85]

◇ Fixed size  $k$  uniform sample from arbitrary size  $N$  stream in one pass

- No need to know stream size in advance
- Include first  $k$  items w.p. 1
- Include item  $n > k$  with probability  $p(n) = k/n, n > k$ 
  - Pick  $j$  uniformly from  $\{1, 2, \dots, n\}$
  - If  $j \leq k$ , swap item  $n$  into location  $j$  in reservoir, discard replaced item

◇ Neat proof shows the uniformity of the sampling method:

- Let  $S_n$  = sample set after  $n$  arrivals



New item: selection probability

$$\text{Prob}[n \in S_n] = p_n := k/n$$

Previously sampled item: induction

$$m \in S_{n-1} \text{ w.p. } p_{n-1} \Rightarrow m \in S_n \text{ w.p. } p_{n-1} * (1 - p_n / k) = p_n$$


## Reservoir Sampling: Skip Counting

- ◇ Simple approach: check each item in turn
  - $O(1)$  per item:
  - Fine if computation time < interarrival time
  - Otherwise build up computation backlog  $O(N)$
- ◇ **Better:** “skip counting”
  - Find random index  $m(n)$  of next selection  $> n$
  - Distribution:  $\text{Prob}[m(n) \leq m] = 1 - (1-p_{n+1}) \cdot (1-p_{n+2}) \cdot \dots \cdot (1-p_m)$
- ◇ Expected number of selections from stream is
$$k + \sum_{k < m \leq N} p_m = k + \sum_{k < m \leq N} k/m = O(k (1 + \ln(N/k)))$$
- ◇ Vitter'85 provided algorithm with this average running time



## Reservoir Sampling via Order Sampling

- ◇ Order sampling a.k.a. bottom-k sample, min-hashing
- ◇ Uniform sampling of stream into reservoir of size  $k$
- ◇ Each arrival  $n$ : generate one-time random value  $r_n \in U[0,1]$ 
  - $r_n$  also known as hash, rank, tag...
- ◇ Store  $k$  items with the smallest random tags



- Each item has same chance of least tag, so uniform
- Fast to implement via priority queue
- Can run on multiple input streams separately, then merge

## Handling Weights

- ◇ So far: uniform sampling from a stream using a reservoir
- ◇ Extend to non-uniform sampling from weighted streams
  - Easy case:  $k=1$
  - Sampling probability  $p(n) = x_n/W_n$  where  $W_n = \sum_{i=1}^n x_i$
- ◇  $k>1$  is harder
  - Can have elements with large weight: would be sampled with prob 1?
- ◇ Number of different weighted order-sampling schemes proposed to realize desired distributional objectives
  - Rank  $r_n = f(u_n, x_n)$  for some function  $f$  and  $u_n \in U[0,1]$
  - k-mins sketches [Cohen 1997], Bottom-k sketches [Cohen Kaplan 2007]
  - [Rosen 1972], Weighted random sampling [Efraimidis Spirakis 2006]
  - Order PPS Sampling [Ohlsson 1990, Rosen 1997]
  - Priority Sampling [Duffield Lund Thorup 2004], [Alon+DLT 2005]



## Weighted random sampling

- ◇ Weighted random sampling [Efraimidis Spirakis 06] generalizes min-wise
  - For each item draw  $r_n$  uniformly at random in range  $[0,1]$
  - Compute the 'tag' of an item as  $r_n^{(1/x_n)}$
  - Keep the items with the  $k$  smallest tags
  - Can prove the correctness of the exponential sampling distribution
- ◇ Can also make efficient via skip counting ideas



## Priority Sampling

- ◇ Each item  $x_i$  given priority  $z_i = x_i / r_i$  with  $r_i$  uniform random in  $(0,1]$
- ◇ Maintain reservoir of  $k+1$  items  $(x_i, z_i)$  of highest priority
- ◇ Estimation
  - Let  $z^* = (k+1)^{\text{st}}$  highest priority
  - Top- $k$  priority items: weight estimate  $x'_i = \max\{x_i, z^*\}$
  - All other items: weight estimate zero
- ◇ Statistics and bounds
  - $x'_i$  unbiased; zero covariance:  $\text{Cov}[x'_i, x'_j] = 0$  for  $i \neq j$
  - Relative variance for any subset sum  $\leq 1/(k-1)$  [Szegedy, 2006]

## Priority Sampling in Databases

### ◇ One Time Sample Preparation

- Compute priorities of all items, sort in decreasing priority order
  - No discard

### ◇ Sample and Estimate

- Estimate any subset sum  $X(S) = \sum_{i \in S} x_i$  by  $X'(S) = \sum_{i \in S} x'_i$  for some  $S' \subset S$
- Method: select items in decreasing priority order

### ◇ Two variants: bounded variance or complexity

1.  $S' =$  first  $k$  items from  $S$ : relative variance bounded  $\leq 1/(k-1)$ 
  - $x'_i = \max\{x_i, z^*\}$  where  $z^* = (k+1)^{\text{st}}$  highest priority in  $S$
2.  $S' =$  items from  $S$  in first  $k$ : execution time  $O(k)$ 
  - $x'_i = \max\{x_i, z^*\}$  where  $z^* = (k+1)^{\text{st}}$  highest priority

[Alon et. al., 2005]

## Making Stream Samples Smarter

- ◇ Observation: we **see** the whole stream, even if we can't store it
  - Can keep more information about sampled items if repeated
  - Simple information: if item sampled, count all repeats
- ◇ Counting Samples [Gibbons & Mattias 98]
  - Sample new items with fixed probability  $p$ , count repeats as  $c_i$
  - Unbiased estimate of total count:  $1/p + (c_i - 1)$
- ◇ Sample and Hold [Estan & Varghese 02]: generalize to weighted keys
  - New key with weight  $b$  sampled with probability  $1 - (1-p)^b$
- ◇ Lower variance compared with independent sampling
  - But sample size will grow as  $pn$
- ◇ Adaptive sample and hold: reduce  $p$  when needed
  - “Sticky sampling”: geometric decreases in  $p$  [Manku, Motwani 02]
  - Much subsequent work tuning decrease in  $p$  to maintain sample size

## Sketch Guided Sampling

- ◇ Go further: avoid sampling the heavy keys as much
  - Uniform sampling will pick from the heavy keys again and again
- ◇ Idea: use an oracle to tell when a key is heavy [Kumar Xu 06]
  - Adjust sampling probability accordingly
- ◇ Can use a “sketch” data structure to play the role of oracle
  - Like a hash table with collisions, tracks approximate frequencies
  - E.g. (Counting) Bloom Filters, Count-Min Sketch
- ◇ Track probability with which key is sampled, use HT estimators
  - Set probability of sampling key with (estimated) weight  $w$  as  $1/(1 + \epsilon w)$  for parameter  $\epsilon$  : decreases as  $w$  increases
  - Decreasing  $\epsilon$  improves accuracy, increases sample size

# Challenges for Smart Stream Sampling

- ◇ Current router constraints
  - Flow tables maintained in fast expensive SRAM
    - To support per packet key lookup at line rate
- ◇ Implementation requirements
  - Sample and Hold: still need per packet lookup
  - Sampled NetFlow: (uniform) sampling reduces lookup rate
    - Easier to implement despite inferior statistical properties
- ◇ Long development times to realize new sampling algorithms
- ◇ Similar concerns affect sampling in other applications
  - Processing large amounts of data needs awareness of hardware
  - Uniform sampling means no coordination needed in distributed setting

# Future for Smarter Stream Sampling

- ◇ Software Defined Networking
  - Current: proprietary software running on special vendor equipment
  - Future: open software and protocols on commodity hardware
- ◇ Potentially offers flexibility in traffic measurement
  - Allocate system resources to measurement tasks as needed
  - Dynamic reconfiguration, fine grained tuning of sampling
  - Stateful packet inspection and sampling for network security
- ◇ Technical challenges:
  - High rate packet processing in software
  - Transparent support from commodity hardware
  - OpenSketch: [\[Yu, Jose, Miao, 2013\]](#)
- ◇ Same issues in other applications: use of commodity programmable HW

# **Stream Sampling: Sampling as Cost Optimization**



## Matching Data to Sampling Analysis

### ♦ Generic problem 1: Counting objects: weight $x_i = 1$

Bernoulli (uniform) sampling with probability  $p$  works fine

- Estimated subset count  $X'(S) = \#\{\text{samples in } S\} / p$
- Relative Variance  $(X'(S)) = (1/p - 1)/X(S)$ 
  - given  $p$ , get any desired accuracy for large enough  $S$




### ♦ Generic problem 2: $x_i$ in Pareto distribution, a.k.a. 80-20 law

- Small proportion of objects possess a large proportion of total weight
  - How to best to sample objects to accurately estimate weight?
- Uniform sampling?
  - likely to omit heavy objects  $\Rightarrow$  big hit on accuracy
  - making selection set  $S$  large doesn't help
- Select  $m$  largest objects ?
  - biased & smaller objects systematically ignored



# Heavy Tails in the Internet and Beyond

- ◇ Files sizes in storage
- ◇ Bytes and packets per network flow
- ◇ Degree distributions in web graph, social networks



## Non-Uniform Sampling

- ◇ Extensive literature: see book by [Tille, “Sampling Algorithms”, 2006]
- ◇ Predates “Big Data”
  - Focus on statistical properties, not so much computational
- ◇ **IPPS**: Inclusion Probability Proportional to Size
  - Variance Optimal for HT Estimation
  - Sampling probabilities for multivariate version: [Chao 1982, Tille 1996]
  - Efficient stream sampling algorithm: [Cohen et. al. 2009]

# Costs of Non-Uniform Sampling

- ◇ Independent sampling from  $n$  objects with weights  $\{x_1, \dots, x_n\}$
- ◇ Goal: find the “best” sampling probabilities  $\{p_1, \dots, p_n\}$
- ◇ **Horvitz-Thompson**: unbiased estimation of each  $x_i$  by


$$x'_i = \begin{cases} x_i/p_i & \text{if weight } i \text{ selected} \\ 0 & \text{otherwise} \end{cases}$$

- ◇ Two costs to balance:
  1. Estimation Variance:  $\text{Var}(x'_i) = x_i^2 (1/p_i - 1)$
  2. Expected Sample Size:  $\sum_i p_i$
- ◇ Minimize Linear Combination Cost:  $\sum_i (x_i^2 (1/p_i - 1) + z^2 p_i)$ 
  - $z$  expresses relative importance of small sample vs. small variance

## Minimal Cost Sampling: IPPS

**IPPS:** Inclusion Probability Proportional to Size

- ◇ Minimize Cost  $\sum_i (x_i^2 (1/p_i - 1) + z^2 p_i)$  subject to  $1 \geq p_i \geq 0$
- ◇ Solution:  $p_i = p_z(x_i) = \min\{1, x_i / z\}$ 
  - small objects ( $x_i < z$ ) selected with probability proportional to size
  - large objects ( $x_i \geq z$ ) selected with probability 1
  - Call  $z$  the “sampling threshold”
  - Unbiased estimator  $x_i/p_i = \max\{x_i, z\}$
- ◇ Perhaps reminiscent of importance sampling, but not the same:
  - make no assumptions concerning distribution of the  $x$



## Error Estimates and Bounds

### ◇ Variance Based:

- HT sampling variance for single object of weight  $x_i$ 
  - $\text{Var}(x'_i) = x_i^2 (1/p_i - 1) = x_i^2 (1/\min\{1, x_i/z\} - 1) \leq z x_i$
- Subset sum  $X(S) = \sum_{i \in S} x_i$  is estimated by  $X'(S) = \sum_{i \in S} x'_i$ 
  - $\text{Var}(X'(S)) \leq z X(S)$

### ◇ Exponential Bounds

- E.g.  $\text{Prob}[X'(S) = 0] \leq \exp(-X(S) / z)$

### ◇ Bounds are simple and powerful

- depend only on subset sum  $X(S)$ , not individual constituents

# Sampled IP Traffic Measurements


- ◇ Packet Sampled NetFlow
    - Sample packet stream in router to limit rate of key lookup: uniform  $1/N$
    - Aggregate sampled packets into flow records by key
  - ◇ Model: packet stream of (key, bytesize) pairs  $\{ (b_i, k_i) \}$
  - ◇ Packet sampled flow record  $(b, k)$  where  $b = \sum \{ b_i : i \text{ sampled} \wedge k_i = k \}$ 
    - HT estimate  $b/N$  of total bytes in flow
  - ◇ Downstream sampling of flow records in measurement infrastructure
    - IPPS sampling, probability  $\min\{1, b/(Nz)\}$
  - ◇ Chained variance bound for any subset sum  $X$  of flows
    - $\text{Var}(X') \leq (z + Nb_{\max}) X$  where  $b_{\max}$  = maximum packet byte size
    - Regardless of how packets are distributed amongst flows
- [Duffield, Lund, Thorup, IEEE ToIT, 2004]

## Estimation Accuracy in Practice

- ◇ Estimate any subset sum comprising at least some fraction  $f$  of weight
- ◇ Suppose: sample size  $m$
- ◇ **Analysis:** typical estimation error  $\epsilon$  (relative standard deviation) obeys

$$\epsilon \leq \frac{1}{\sqrt{f m}}$$

Estimate fraction  $f = 0.1\%$   
with typical relative error  
12%:




- ◇  $2^{16}$  = same storage needed for aggregates over 16 bit address prefixes
  - But sampling gives more flexibility to estimate traffic within aggregates



## Heavy Hitters: Exact vs. Aggregate vs. Sampled

- ◇ Sampling does not tell you where the interesting features are
  - But does speed up the ability to find them with existing tools
- ◇ Example: Heavy Hitter Detection
  - Setting: Flow records reporting 10GB/s traffic stream
  - Aim: find Heavy Hitters = IP prefixes comprising  $\geq 0.1\%$  of traffic
  - Response time needed: 5 minute
- ◇ Compare:
  - Exact: 10GB/s x 5 minutes yields upwards of 300M flow records
  - 64k aggregates over 16 bit prefixes: no deeper drill-down possible
  - Sampled: 64k flow records: **any** aggregate  $\geq 0.1\%$  accurate to 10%



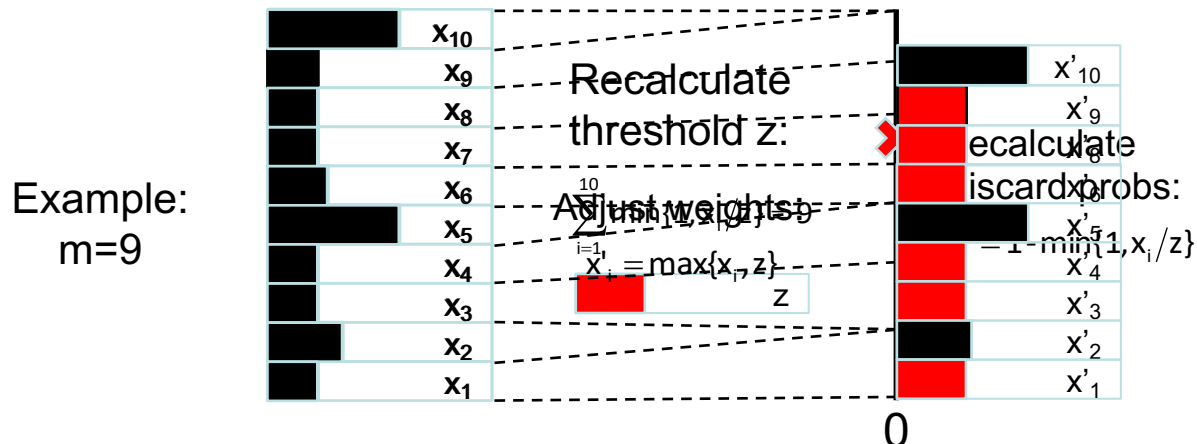
## Cost Optimization for Sampling

Several different approaches optimize for different objectives:

1. **Fixed Sample Size IPPS Sample**
  - Variance Optimal sampling: minimal variance unbiased estimation
2. **Structure Aware Sampling**
  - Improve estimation accuracy for subnet queries using topological cost
3. **Fair Sampling**
  - Adaptively balance sampling budget over subpopulations of flows
  - Uniform estimation accuracy regardless of subpopulation size
4. **Stable Sampling**
  - Increase stability of sample set by imposing cost on changes


## IPPS Stream Reservoir Sampling

- ◇ Each arriving item:
  - Provisionally include item in reservoir
  - If  $m+1$  items, discard 1 item randomly
    - Calculate threshold  $z$  to sample  $m$  items on average:  $z$  solves  $\sum_i p_z(x_i) = m$
    - Discard item  $i$  with probability  $q_i = 1 - p_z(x_i)$
    - Adjust  $m$  surviving  $x_i$  with Horvitz-Thompson  $x'_i = x_i / p_i = \max\{x_i, z\}$
- ◇ Efficient Implementation:
  - Computational cost  $O(\log m)$  per item, amortized cost  $O(\log \log m)$



# Structure (Un)Aware Sampling

- ◇ Sampling is oblivious to structure in keys (IP address hierarchy)
  - Estimation disperses the weight of discarded items to surviving samples




- ◇ Queries structure aware: subset sums over related keys (IP subnets)
  - Accuracy on LHS is decreased by discarding weight on RHS

## Localizing Weight Redistribution

- ◇ Initial weight set  $\{x_i : i \in S\}$  for some  $S \subset \Omega$ 
  - E.g.  $\Omega$  = possible IP addresses,  $S$  = observed IP addresses
- ◇ Attribute “range cost”  $C(\{x_i : i \in R\})$  for each weight subset  $R \subseteq S$ 
  - Possible factors for Range Cost:
    - Sampling variance
    - Topology e.g. height of lowest common ancestor
  - Heuristics:  $R^* =$  Nearest Neighbor  $\{x_i, x_j\}$  of minimal  $x_i x_j$
- ◇ Sample  $k$  items from  $S$ :
  - Progressively remove one item from subset with minimal range cost:
  - While( $|S| > k$ )
    - Find  $R^* \subseteq S$  of minimal range cost.
    - Remove a weight from  $R^*$  w/ VarOpt


[Cohen, Cormode, Duffield; PVLDB 2011]



Order of magnitude reduction in average subnet error vs. VarOpt

## Fair Sampling Across Subpopulations

- ◇ Analysis queries often focus on specific subpopulations
  - E.g. networking: different customers, user applications, network paths
- ◇ Wide variation in subpopulation size
  - 5 orders of magnitude variation in traffic on interfaces of access router
- ◇ If uniform sampling across subpopulations:
  - Poor estimation accuracy on subset sums within small subpopulations



Color = subpopulation

▲ , ▲ = interesting items

– occurrence proportional to subpopulation size

Uniform Sampling across subpopulations:

– Difficult to track proportion of interesting items within small subpopulations: ▲

## Fair Sampling Across Subpopulations

- ◇ Minimize **relative** variance by sharing budget  $m$  over subpopulations
  - Total  $n$  objects in subpopulations  $n_1, \dots, n_d$  with  $\sum_i n_i = n$
  - Allocate budget  $m_i$  to each subpopulation  $n_i$  with  $\sum_i m_i = m$
- ◇ Minimize average population relative variance  $R = \text{const.} \sum_i 1/m_i$
- ◇ Theorem:
  - $R$  minimized when  $\{m_i\}$  are Max-Min Fair share of  $m$  under demands  $\{n_i\}$
- ◇ Streaming
  - Problem: don't know subpopulation sizes  $\{n_i\}$  in advance
- ◇ Solution: progressive fair sharing as reservoir sample
  - Provisionally include each arrival
  - Discard 1 item as VarOpt sample from any maximal subpopulation
- ◇ Theorem [Duffield; Sigmetrics 2012]:
  - Max-Min Fair at all times; equality in distribution with VarOpt samples  $\{m_i$  from  $n_i\}$

## Stable Sampling




- ◇ **Setting**: Sampling a population over successive periods
- ◇ Sample independently at each time period?
  - Cost associated with sample churn
  - Time series analysis of set of relatively stable keys
- ◇ Find sampling probabilities through cost minimization
  - Minimize Cost = Estimation Variance +  $z * E[\text{\#Churn}]$
- ◇ Size  $m$  sample with maximal expected churn  $D$ 
  - weights  $\{x_i\}$ , previous sampling probabilities  $\{p_i\}$
  - find new sampling probabilities  $\{q_i\}$  to minimize cost of taking  $m$  samples
  - Minimize  $\sum_i x_i^2 / q_i$  subject to  $1 \geq q_i \geq 0$ ,  $\sum_i q_i = m$  and  $\sum_i |p_i - q_i| \leq D$

[Cohen, Cormode, Duffield, Lund 13]



## Summary of Part 1

- ◇ Sampling as a powerful, general summarization technique
- ◇ Unbiased estimation via Horvitz-Thompson estimators
- ◇ Sampling from streams of data
  - Uniform sampling: reservoir sampling
  - Weighted generalizations: sample and hold, counting samples
- ◇ Advances in stream sampling
  - The cost principle for sample design, and IPPS methods
  - Threshold, priority and VarOpt sampling
  - Extending the cost principle:
    - structure aware, fair sampling, stable sampling, sketch guided



# Sampling for Big Data

Graham Cormode, University of Warwick

[G.Cormode@warwick.ac.uk](mailto:G.Cormode@warwick.ac.uk)

Nick Duffield, Texas A&M University

[Nick.Duffield@gmail.com](mailto:Nick.Duffield@gmail.com)

